

Bernoulli example

■ Define the working directory and load CmdStan.m

```
In[36]:= (* Linux *)
SetDirectory["~/GitHub/MathematicaStan/Examples/Bernoulli/"]

(* Windows *)
(* SetDirectory["C:\\Users\\USER_NAME\\Documents\\Mathematica\\STAN\\Examples\\Bernoulli"] *)

Needs["CmdStan`"]

Out[36]= /IS006139/home/pix/GitHub/MathematicaStan/Examples/Bernoulli
```

■ Generate the Bernoulli Stan code and compile it

```
In[3]:= stanCode="data {
  int<lower=0> N;
  int<lower=0,upper=1> y[N];
}
parameters {
  real<lower=0,upper=1> theta;
}
model {
  theta ~ beta(1,1);
  for (n in 1:N)
    y[n] ~ bernoulli(theta);
}";
StanCodeExport["bernoulli",stanCode]

(* Compile your code.
 * Caveat: this can take some time
 *)
StanCompile["bernoulli"]
```

```
Out[4]= bernoulli.stan
```

```
Out[5]= make: '/IS006139/home/pix/GitHub/MathematicaStan/Examples/Bernoulli/bernoulli' is up to date.
```

■ Generate some data and save them (RDump file)

```
In[6]:= n=5000;
y=Table[Random[BernoulliDistribution[0.2016]],{i,1,n}];

RDumpExport["bernoulli",{{"N",n},{ "Y",y}}];
```

■ Run Stan and get result

■ Use sample method (ONE job)

```
In[9]:= StanRunSample["bernoulli"]

output=StanImport["output.csv"];

Out[9]= method = sample (Default)
sample
num_samples = 1000 (Default)
```

```

num_warmup = 1000 (Default)
save_warmup = 0 (Default)
thin = 1 (Default)
adapt
  engaged = 1 (Default)
  gamma = 0.050000000000000003 (Default)
  delta = 0.800000000000000004 (Default)
  kappa = 0.75 (Default)
  t0 = 10 (Default)
  init_buffer = 75 (Default)
  term_buffer = 50 (Default)
  window = 25 (Default)
algorithm = hmc (Default)
  hmc
    engine = nuts (Default)
      nuts
        max_depth = 10 (Default)
        metric = diag_e (Default)
        stepsize = 1 (Default)
        stepsize_jitter = 0 (Default)
id = 0 (Default)
data
  file = /IS006139/home/pix/GitHub/MathematicaStan/Examples/Bernoulli/bernoulli.data.R
init = 2 (Default)
random
  seed = 3919364410
output
  file = /IS006139/home/pix/GitHub/MathematicaStan/Examples/Bernoulli/output.csv
  diagnostic_file = (Default)
  refresh = 100 (Default)

```

Gradient evaluation took 0.000399 seconds
 1000 transitions using 10 leapfrog steps per transition would take 3.99 seconds.
 Adjust your expectations accordingly!

```

Iteration: 1 / 2000 [ 0%] (Warmup)
Iteration: 100 / 2000 [ 5%] (Warmup)
Iteration: 200 / 2000 [ 10%] (Warmup)
Iteration: 300 / 2000 [ 15%] (Warmup)
Iteration: 400 / 2000 [ 20%] (Warmup)
Iteration: 500 / 2000 [ 25%] (Warmup)
Iteration: 600 / 2000 [ 30%] (Warmup)
Iteration: 700 / 2000 [ 35%] (Warmup)
Iteration: 800 / 2000 [ 40%] (Warmup)
Iteration: 900 / 2000 [ 45%] (Warmup)
Iteration: 1000 / 2000 [ 50%] (Warmup)
Iteration: 1001 / 2000 [ 50%] (Sampling)
Iteration: 1100 / 2000 [ 55%] (Sampling)
Iteration: 1200 / 2000 [ 60%] (Sampling)
Iteration: 1300 / 2000 [ 65%] (Sampling)
Iteration: 1400 / 2000 [ 70%] (Sampling)
Iteration: 1500 / 2000 [ 75%] (Sampling)
Iteration: 1600 / 2000 [ 80%] (Sampling)
Iteration: 1700 / 2000 [ 85%] (Sampling)
Iteration: 1800 / 2000 [ 90%] (Sampling)

```

```
Iteration: 1900 / 2000 [ 95%] (Sampling)
Iteration: 2000 / 2000 [100%] (Sampling)
```

```
Elapsed Time: 1.47087 seconds (Warm-up)
              1.7528 seconds (Sampling)
              3.22366 seconds (Total)
```

■ Use the results

■ List Header

```
In[11]:= StanImportHeader[output]
```

```
Out[11]= {{lp__, 1}, {accept_stat__, 2}, {stepsize__, 3}, {treedepth__, 4},
          {n_leapfrog__, 5}, {divergent__, 6}, {energy__, 7}, {theta, 8}}
```

■ Show sample matrix

```
In[12]:= Dimensions[StanImportData[output]]
```



```
Take[StanImportData[output],3]
```

```
Out[12]= {1000, 8}
```

```
Out[13]= {{-2508.15, 0.574074, 1.5896, 1., 1., 0., 2509.9, 0.203886},
          {-2508.34, 0.948844, 1.5896, 1., 3., 0., 2508.41, 0.196025},
          {-2508.47, 0.969459, 1.5896, 1., 3., 0., 2508.61, 0.206291}}
```

■ Plot θ sample and histogram

```
In[14]:= ListLinePlot[Flatten[StanVariableColumn["theta",output]],PlotLabel->" $\theta$ "]  
Histogram[Flatten[StanVariableColumn["theta",output]],PlotLabel->" $\theta$ "]
```



■ Maximize likelihood with StanRunOptimize

```
In[16]:= StanRunOptimize["bernoulli"]
```

```
Out[16]= method = optimize
          optimize
            algorithm = lbfgs (Default)
            lbfgs
              init_alpha = 0.001 (Default)
              tol_obj = 9.999999999999998e-13 (Default)
              tol_rel_obj = 10000 (Default)
              tol_grad = 1e-08 (Default)
              tol_rel_grad = 10000000 (Default)
              tol_param = 1e-08 (Default)
              history_size = 5 (Default)
            iter = 2000 (Default)
            save_iterations = 0 (Default)
          id = 0 (Default)
        data
          file = /IS006139/home/pix/GitHub/MathematicaStan/Examples/Bernoulli/bernoulli.data.R
        init = 2 (Default)
      random
        seed = 3919368427
    output
      file = /IS006139/home/pix/GitHub/MathematicaStan/Examples/Bernoulli/output.csv
      diagnostic_file = (Default)
      refresh = 100 (Default)

initial log joint probability = -2513.39
  Iter      log prob      ||dx||      ||grad||      alpha      alpha0  # evals  Notes
    3      -2506.17      0.00131072      0.0238056      0.9687      0.9687      4
Optimization terminated normally:
  Convergence detected: relative gradient magnitude is below tolerance
```

■ Options manipulation

```
In[17]:= StanSetOptionOptimize["output.file","output_optimize.csv"];
StanSetOptionOptimize["method.optimize.iter",100];
StanSetOptionOptimize["method.optimize.algorithm","bfgs"];
StanSetOptionOptimize["method.optimize.algorithm.bfgs.tol_grad",10.^-5];
StanOptionOptimize[]

(* re-run the solver with the new options *)
StanRunOptimize["bernoulli"]

Out[21]= {{method.optimize.algorithm.bfgs.tol_grad, 0.00001}, {method.optimize.algorithm, bfgs},
{method.optimize.iter, 100}, {output.file, output_optimize.csv}}

Out[22]= method = optimize
optimize
algorithm = bfgs
bfgs
init_alpha = 0.001 (Default)
tol_obj = 9.999999999999998e-13 (Default)
tol_rel_obj = 10000 (Default)
tol_grad = 1.0000000000000001e-05
tol_rel_grad = 10000000 (Default)
tol_param = 1e-08 (Default)
iter = 100
save_iterations = 0 (Default)
id = 0 (Default)
data
file = /IS006139/home/pix/GitHub/MathematicaStan/Examples/Bernoulli/bernoulli.data.R
init = 2 (Default)
random
seed = 3919368500
output
file = output_optimize.csv
diagnostic_file = (Default)
refresh = 100 (Default)

initial log joint probability = -2517.89
  Iter      log prob      ||dx||      ||grad||      alpha      alpha0  # evals  Notes
    3      -2506.17    0.00210032    0.0535299    0.9571    0.9571      4
Optimization terminated normally:
Convergence detected: relative gradient magnitude is below tolerance
```

■ Overwrite and/or reset option

```
In[23]:= StanOptionOptimize[]
StanSetOptionOptimize["method.optimize.iter",2016];
StanOptionOptimize[]

Out[23]= {{method.optimize.algorithm.bfgs.tol_grad, 0.00001}, {method.optimize.algorithm, bfgs},
{method.optimize.iter, 100}, {output.file, output_optimize.csv}}

Out[25]= {{method.optimize.algorithm.bfgs.tol_grad, 0.00001}, {method.optimize.algorithm, bfgs},
{method.optimize.iter, 2016}, {output.file, output_optimize.csv}}
```

■ Remove all method* options

```
In[26]:= StanOptionOptimize[]
StanRemoveOptionOptimize["method*"];
StanOptionOptimize[]

Out[26]:= {{method.optimize.algorithm.bfgs.tol_grad, 0.00001}, {method.optimize.algorithm, bfgs},
           {method.optimize.iter, 2016}, {output.file, output_optimize.csv}}

Out[28]:= {{output.file, output_optimize.csv}}
```

■ Erase all options

```
In[29]:= StanOptionOptimize[]
StanResetOptionOptimize[];
StanOptionOptimize[]

Out[29]:= {{output.file, output_optimize.csv}}

Out[31]:= {}
```

■ Parallel Sampling

■ Redo the previous computation with 4 jobs in parallel (ONLY works under Linux for the moment)

```
In[32]:= StanRunSample["bernoulli",4] (* 4 jobs *)
output=StanImport["output.csv"];
ListLinePlot[Flatten[StanVariableColumn["theta",output]],PlotLabel->"θ"]
Histogram[Flatten[StanVariableColumn["theta",output]],PlotLabel->"θ"]

Out[32]= method = sample (Default)
  sample
    num_samples = 1000 (Default)
    num_warmup = 1000 (Default)
    save_warmup = 0 (Default)
    thin = 1 (Default)
  adapt
    engaged = 1 (Default)
    gamma = 0.050000000000000003 (Default)
    delta = 0.80000000000000004 (Default)
    kappa = 0.75 (Default)
    t0 = 10 (Default)
    init_buffer = 75 (Default)
    term_buffer = 50 (Default)
    window = 25 (Default)
  algorithm = hmc (Default)
  hmc
    engine = nuts (Default)
    nuts
      max_depth = 10 (Default)
      metric = diag_e (Default)
      stepsize = 1 (Default)
      stepsize_jitter = 0 (Default)
  id = 2
  data
    file = /IS006139/home/pix/GitHub/MathematicaStan/Examples/Bernoulli/bernoulli.data.R
  init = 2 (Default)
  random
    seed = 3919368663
  output
    file = /IS006139/home/pix/GitHub/MathematicaStan/Examples/Bernoulli/output_2.csv
```

```
diagnostic_file = (Default)
refresh = 100 (Default)
```

Gradient evaluation took 0.000377 seconds
 1000 transitions using 10 leapfrog steps per transition would take 3.77 seconds.
 Adjust your expectations accordingly!

```
method = sample (Default)
sample
  num_samples = 1000 (Default)
  num_warmup = 1000 (Default)
  save_warmup = 0 (Default)
  thin = 1 (Default)
  adapt
    engaged = 1 (Default)
    gamma = 0.050000000000000003 (Default)
    delta = 0.80000000000000004 (Default)
    kappa = 0.75 (Default)
    t0 = 10 (Default)
    init_buffer = 75 (Default)
    term_buffer = 50 (Default)
    window = 25 (Default)
  algorithm = hmc (Default)
  hmc
    engine = nuts (Default)
    nuts
      max_depth = 10 (Default)
      metric = diag_e (Default)
      stepsize = 1 (Default)
      stepsize_jitter = 0 (Default)
id = 4
data
  file = /IS006139/home/pix/GitHub/MathematicaStan/Examples/Bernoulli/bernoulli.data.R
init = 2 (Default)
random
  seed = 3919368665
output
  file = /IS006139/home/pix/GitHub/MathematicaStan/Examples/Bernoulli/output_4.csv
  diagnostic_file = (Default)
  refresh = 100 (Default)

method = sample (Default)
sample
  num_samples = 1000 (Default)
  num_warmup = 1000 (Default)
  save_warmup = 0 (Default)
  thin = 1 (Default)
  adapt
    engaged = 1 (Default)
    gamma = 0.050000000000000003 (Default)
method = sample (Default)
sample
  delta = 0.80000000000000004 (Default)
  num_samples = 1000 (Default)
  num_warmup = 1000 (Default)
```



```

    kappa = 0.75 (Default)
save_warmup = 0 (Default)
thin = 1 (Default)
adapt
  engaged = 1 (Default)
  t0 = 10 (Default)
  init_buffer = 75 (Default)
  term_buffer = 50 (Default)
  window = 25 (Default)
  gamma = 0.050000000000000003 (Default)
algorithm = hmc (Default)
  hmc
    engine = nuts (Default)
    nuts
      delta = 0.80000000000000004 (Default)
      max_depth = 10 (Default)
      metric = diag_e (Default)
      kappa = 0.75 (Default)
      stepsize = 1 (Default)
      t0 = 10 (Default)
      init_buffer = 75 (Default)
      stepsize_jitter = 0 (Default)
      term_buffer = 50 (Default)
      window = 25 (Default)
id = 3
data
  algorithm = hmc (Default)
  hmc
    file = /IS006139/home/pix/GitHub/MathematicaStan/Examples/Bernoulli/bernoulli.data.R
    engine = nuts (Default)
init = 2 (Default)
  nuts
random
  max_depth = 10 (Default)
  seed = 3919368665
output
  metric = diag_e (Default)
  file = /IS006139/home/pix/GitHub/MathematicaStan/Examples/Bernoulli/output_3.csv
  diagnostic_file = (Default)
  refresh = 100 (Default)
  stepsize = 1 (Default)

  stepsize_jitter = 0 (Default)
id = 1
data
  file = /IS006139/home/pix/GitHub/MathematicaStan/Examples/Bernoulli/bernoulli.data.R
init = 2 (Default)
random
  seed = 3919368665
output
  file = /IS006139/home/pix/GitHub/MathematicaStan/Examples/Bernoulli/output_1.csv
  diagnostic_file = (Default)
  refresh = 100 (Default)

```

Gradient evaluation took 0.000555 seconds

1000 transitions using 10 leapfrog steps per transition would take 5.55 seconds.

Adjust your expectations accordingly!

Gradient evaluation took 0.000635 seconds

Gradient evaluation took 0.000633 seconds

1000 transitions using 10 leapfrog steps per transition would take 6.35 seconds.

1000 transitions using 10 leapfrog steps per transition would take 6.33 seconds.

Adjust your expectations accordingly!

Adjust your expectations accordingly!

```
Iteration: 1 / 2000 [ 0%] (Warmup)
Iteration: 1 / 2000 [ 0%] (Warmup)
Iteration: 1 / 2000 [ 0%] (Warmup)
Iteration: 1 / 2000 [ 0%] (Warmup)
Iteration: 100 / 2000 [ 5%] (Warmup)
Iteration: 100 / 2000 [ 5%] (Warmup)
Iteration: 100 / 2000 [ 5%] (Warmup)
Iteration: 100 / 2000 [ 5%] (Warmup)
Iteration: 200 / 2000 [ 10%] (Warmup)
Iteration: 200 / 2000 [ 10%] (Warmup)
Iteration: 200 / 2000 [ 10%] (Warmup)
Iteration: 200 / 2000 [ 10%] (Warmup)
Iteration: 300 / 2000 [ 15%] (Warmup)
Iteration: 300 / 2000 [ 15%] (Warmup)
Iteration: 300 / 2000 [ 15%] (Warmup)
Iteration: 300 / 2000 [ 15%] (Warmup)
Iteration: 400 / 2000 [ 20%] (Warmup)
Iteration: 400 / 2000 [ 20%] (Warmup)
Iteration: 400 / 2000 [ 20%] (Warmup)
Iteration: 400 / 2000 [ 20%] (Warmup)
Iteration: 500 / 2000 [ 25%] (Warmup)
Iteration: 500 / 2000 [ 25%] (Warmup)
Iteration: 500 / 2000 [ 25%] (Warmup)
Iteration: 500 / 2000 [ 25%] (Warmup)
Iteration: 600 / 2000 [ 30%] (Warmup)
Iteration: 600 / 2000 [ 30%] (Warmup)
Iteration: 600 / 2000 [ 30%] (Warmup)
Iteration: 600 / 2000 [ 30%] (Warmup)
Iteration: 700 / 2000 [ 35%] (Warmup)
Iteration: 700 / 2000 [ 35%] (Warmup)
Iteration: 700 / 2000 [ 35%] (Warmup)
Iteration: 700 / 2000 [ 35%] (Warmup)
Iteration: 800 / 2000 [ 40%] (Warmup)
Iteration: 800 / 2000 [ 40%] (Warmup)
Iteration: 800 / 2000 [ 40%] (Warmup)
Iteration: 800 / 2000 [ 40%] (Warmup)
Iteration: 900 / 2000 [ 45%] (Warmup)
Iteration: 900 / 2000 [ 45%] (Warmup)
Iteration: 900 / 2000 [ 45%] (Warmup)
Iteration: 900 / 2000 [ 45%] (Warmup)
Iteration: 1000 / 2000 [ 50%] (Warmup)
Iteration: 1001 / 2000 [ 50%] (Sampling)
```

```

Iteration: 1000 / 2000 [ 50%] (Warmup)
Iteration: 1001 / 2000 [ 50%] (Sampling)
Iteration: 1000 / 2000 [ 50%] (Warmup)
Iteration: 1000 / 2000 [ 50%] (Warmup)
Iteration: 1001 / 2000 [ 50%] (Sampling)
Iteration: 1001 / 2000 [ 50%] (Sampling)
Iteration: 1100 / 2000 [ 55%] (Sampling)
Iteration: 1100 / 2000 [ 55%] (Sampling)
Iteration: 1100 / 2000 [ 55%] (Sampling)
Iteration: 1100 / 2000 [ 55%] (Sampling)
Iteration: 1200 / 2000 [ 60%] (Sampling)
Iteration: 1200 / 2000 [ 60%] (Sampling)
Iteration: 1200 / 2000 [ 60%] (Sampling)
Iteration: 1200 / 2000 [ 60%] (Sampling)
Iteration: 1300 / 2000 [ 65%] (Sampling)
Iteration: 1300 / 2000 [ 65%] (Sampling)
Iteration: 1300 / 2000 [ 65%] (Sampling)
Iteration: 1400 / 2000 [ 70%] (Sampling)
Iteration: 1300 / 2000 [ 65%] (Sampling)
Iteration: 1400 / 2000 [ 70%] (Sampling)
Iteration: 1400 / 2000 [ 70%] (Sampling)
Iteration: 1500 / 2000 [ 75%] (Sampling)
Iteration: 1400 / 2000 [ 70%] (Sampling)
Iteration: 1500 / 2000 [ 75%] (Sampling)
Iteration: 1600 / 2000 [ 80%] (Sampling)
Iteration: 1500 / 2000 [ 75%] (Sampling)
Iteration: 1500 / 2000 [ 75%] (Sampling)
Iteration: 1600 / 2000 [ 80%] (Sampling)
Iteration: 1700 / 2000 [ 85%] (Sampling)
Iteration: 1600 / 2000 [ 80%] (Sampling)
Iteration: 1600 / 2000 [ 80%] (Sampling)
Iteration: 1800 / 2000 [ 90%] (Sampling)
Iteration: 1700 / 2000 [ 85%] (Sampling)
Iteration: 1700 / 2000 [ 85%] (Sampling)
Iteration: 1700 / 2000 [ 85%] (Sampling)
Iteration: 1900 / 2000 [ 95%] (Sampling)
Iteration: 1800 / 2000 [ 90%] (Sampling)
Iteration: 1800 / 2000 [ 90%] (Sampling)
Iteration: 2000 / 2000 [100%] (Sampling)

```

```

Elapsed Time: 1.66671 seconds (Warm-up)
              1.50367 seconds (Sampling)
              3.17038 seconds (Total)

```

```

Iteration: 1800 / 2000 [ 90%] (Sampling)
Iteration: 1900 / 2000 [ 95%] (Sampling)
Iteration: 2000 / 2000 [100%] (Sampling)

```

```

Elapsed Time: 1.72077 seconds (Warm-up)
              1.60987 seconds (Sampling)
              3.33064 seconds (Total)

```

```

Iteration: 1900 / 2000 [ 95%] (Sampling)
Iteration: 1900 / 2000 [ 95%] (Sampling)
Iteration: 2000 / 2000 [100%] (Sampling)

```

```

Elapsed Time: 1.7269 seconds (Warm-up)

```

1.80656 seconds (Sampling)
3.53347 seconds (Total)

Iteration: 2000 / 2000 [100%] (Sampling)

Elapsed Time: 1.69561 seconds (Warm-up)
1.82515 seconds (Sampling)
3.52076 seconds (Total)

