

Horse shoe prior example

This example reproduces some results of <https://ariddell.org/horseshoe-prior-with-stan.html>

■ Define the working directory and load CmdStan.m

```
In[28]:= (* Linux *)
SetDirectory["~/GitHub/MathematicaStan/Examples/HorseShoePrior"]

(* Windows *)
(* SetDirectory["C:\\Users\\USER_NAME\\Documents\\Mathematica\\STAN\\Examples\\HorseShoePrior"]; *)

Needs["CmdStan`"]

Out[28]= /IS006139/home/pix/GitHub/MathematicaStan/Examples/HorseShoePrior
```

■ Generate the Horse Shoe Prior Stan code and compile it

```
In[3]:= stanCode="data {
    int<lower=0> n;
    int<lower=0> p;
    matrix[n,p] X;
    vector[n] y;
}
parameters {
    vector[p] beta;
    vector<lower=0>[p] lambda;
    real<lower=0> tau;
    real<lower=0> sigma;
}
model {
    lambda ~ cauchy(0, 1);
    tau ~ cauchy(0, 1);
    for (i in 1:p)
        beta[i] ~ normal(0, lambda[i] * tau);
    y ~ normal(X * beta, sigma);
}";
StanCodeExport["horseShoePrior",stanCode]

StanCompile["horseShoePrior"]

Out[4]= horseShoePrior.stan

Out[5]= make:
'/IS006139/home/pix/GitHub/MathematicaStan/Examples/HorseShoePrior/horseShoePrior' is up to date.
```

■ Load data and save them (RDump file)

```
In[6]:= yTest=Import["./y-test.dat","List"];
yTrain=Import["./y-train.dat","List"];
XTest= Import["./X-test.dat","Table"];
XTrain= Import["./X-train.dat","Table"];

(* betaLabel is only used for plot legends *)
betaLabel=StringSplit["age sex bmi map tc ldl hdl tch ltg glu age^2 bmi^2 map^2 tc^2 ldl^2 hdl^2 tch^2"]

(* Export data *)
RDumpExport["horseShoePrior",{{"n",Dimensions[XTrain][[1]]},{ "p",Dimensions[XTrain][[2]]},{ "X",XTrain}},{"X",XTrain}]

(* Here we just perform an Ordinary Least Squares to get the residue value *)
betaOLS=LeastSquares[XTrain,yTrain];
Norm[XTest.betaOLS-yTest]^2/Length[yTest]
Norm[Mean[yTrain]-yTest]^2/Length[yTest]

Out[13]= 0.670749
Out[14]= 0.965737
```

■ Run Stan and get result

```
In[15]:= (* use the same seed as the blog post *)
StanSetOptionSample["random seed",5]

Out[15]= {{random seed, 5}}

In[16]:= (* Run stan *)
StanRunSample["horseShoePrior"]

Out[16]= method = sample (Default)
  sample
    num_samples = 1000 (Default)
    num_warmup = 1000 (Default)
    save_warmup = 0 (Default)
    thin = 1 (Default)
  adapt
    engaged = 1 (Default)
    gamma = 0.050000000000000003 (Default)
    delta = 0.80000000000000004 (Default)
    kappa = 0.75 (Default)
    t0 = 10 (Default)
    init_buffer = 75 (Default)
    term_buffer = 50 (Default)
    window = 25 (Default)
  algorithm = hmc (Default)
  hmc
    engine = nuts (Default)
    nuts
      max_depth = 10 (Default)
      metric = diag_e (Default)
      stepsize = 1 (Default)
      stepsize_jitter = 0 (Default)
  id = 0 (Default)
  data
    file = /IS006139/home/pix/GitHub/MathematicaStan/Examples/HorseShoePrior/horseShoePrior.data.R
  init = 2 (Default)
```

```

random
  seed = 5
output
  file = /IS006139/home/pix/GitHub/MathematicaStan/Examples/HorseShoePrior/output.csv
  diagnostic_file = (Default)
  refresh = 100 (Default)

```

Gradient evaluation took 0.000187 seconds
 1000 transitions using 10 leapfrog steps per transition would take 1.87 seconds.
 Adjust your expectations accordingly!

```

Iteration:   1 / 2000 [ 0%] (Warmup)
Iteration:  100 / 2000 [ 5%] (Warmup)
Iteration:  200 / 2000 [10%] (Warmup)
Iteration:  300 / 2000 [15%] (Warmup)
Iteration:  400 / 2000 [20%] (Warmup)
Iteration:  500 / 2000 [25%] (Warmup)
Iteration:  600 / 2000 [30%] (Warmup)
Iteration:  700 / 2000 [35%] (Warmup)
Iteration:  800 / 2000 [40%] (Warmup)
Iteration:  900 / 2000 [45%] (Warmup)
Iteration: 1000 / 2000 [50%] (Warmup)
Iteration: 1001 / 2000 [50%] (Sampling)
Iteration: 1100 / 2000 [55%] (Sampling)
Iteration: 1200 / 2000 [60%] (Sampling)
Iteration: 1300 / 2000 [65%] (Sampling)
Iteration: 1400 / 2000 [70%] (Sampling)
Iteration: 1500 / 2000 [75%] (Sampling)
Iteration: 1600 / 2000 [80%] (Sampling)
Iteration: 1700 / 2000 [85%] (Sampling)
Iteration: 1800 / 2000 [90%] (Sampling)
Iteration: 1900 / 2000 [95%] (Sampling)
Iteration: 2000 / 2000 [100%] (Sampling)

```

```

Elapsed Time: 84.1096 seconds (Warm-up)
              34.6386 seconds (Sampling)
              118.748 seconds (Total)

```

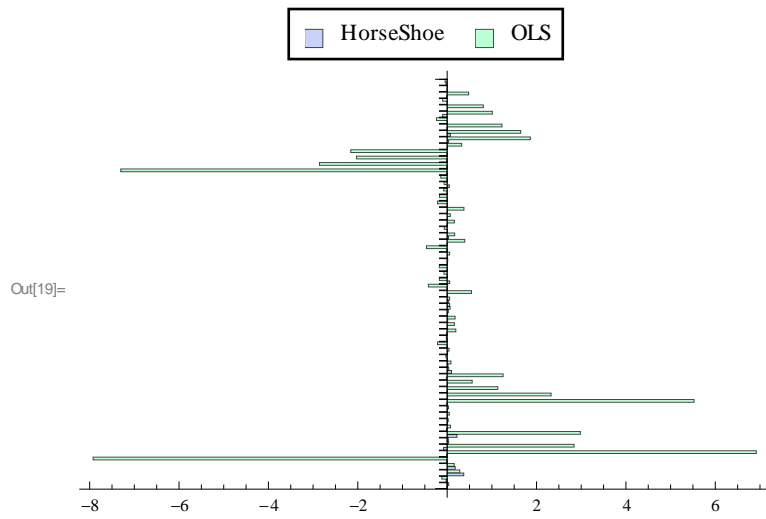
■ Use the results

```
In[17]:= output = StanImport["output.csv"];
```

Compute beta mean and compare it to OLS solution

```
In[18]:= beta = Mean[StanVariableColumn["beta", output]];
```

```
In[19]:= BarChart[Transpose[{beta,betaOLS}],ChartLegends->Placed[{"HorseShoe","OLS"},Top],BarOrigin->Left]
```



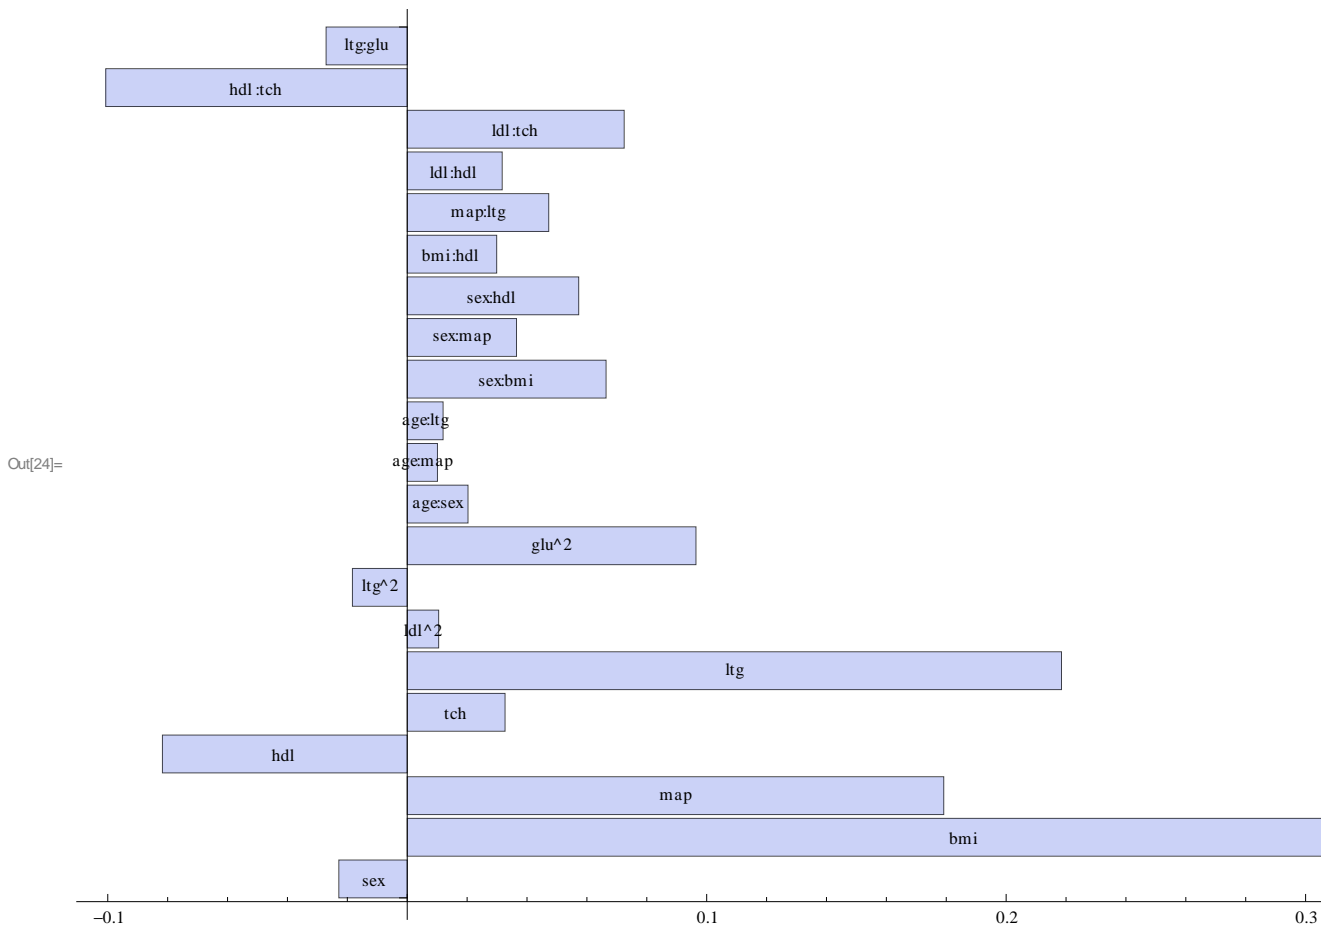
■ Now select all beta such that $|\text{beta}| > 0.01$

```
In[20]:= selectedBeta=Map[ (Abs[#]>0.01)&,&beta] ;
Print["The ",Count[selectedBeta,True]," variables are:\n",
(betaLabel[[selectedBeta=Flatten[Position[selectedBeta,True]]]])];

selectedBetaLabel=betaLabel[[selectedBeta]];
selectedBeta=beta[[selectedBeta]];

BarChart[selectedBeta,ChartLabels->Placed[selectedBetaLabel,Center],BarOrigin->Left]

The 21 variables are:
{sex, bmi, map, hdl, tch, ltg, ldl^2, ltg^2, glu^2, age:sex, age:map, age:ltg,
sex:bmi, sex:map, sex:hdl, bmi:hdl, map:ltg, ldl:hdl, ldl:tch, hdl:tch, ltg:glu}
```



■ It is interesting to notice that the pruned beta has a better generalization than the OLS solution:

```
In[25]:= prunedBeta=Map[ If[Abs[#]>0.01,&#,0]&,&beta] ;
Norm[XTest.betaOLS-yTest]^2/Length[yTest]
Norm[XTest.prunedBeta-yTest]^2/Length[yTest]
```

Out[26]= 0.670749

Out[27]= 0.50374